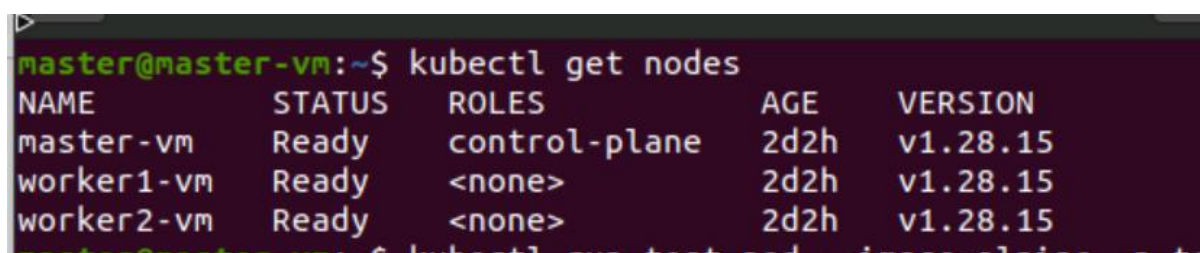


Kubernetes Multi-Tenant Project

Step 1: Check if Any Worker Node is Ready

Run the following command to check the status of worker nodes:

```
kubectl get nodes
```



```
master@master-vm:~$ kubectl get nodes
NAME              STATUS    ROLES    AGE   VERSION
master-vm         Ready    control-plane   2d2h   v1.28.15
worker1-vm        Ready    <none>         2d2h   v1.28.15
worker2-vm        Ready    <none>         2d2h   v1.28.15
```

Step 2: Install Calico for Networking

Apply the Calico manifest to enable networking:

```
kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml
```

Step 3: Create Namespaces for Tenants

To isolate tenants, create separate namespaces:

```
kubectl create namespace tenant-a
```

```
kubectl create namespace tenant-b
```

Step 4: Create Folder Structure for YAML Files

Create the folder structure to organize YAML files for each tenant:

```
mkdir -p ~/k8s-multi-tenant/tenant-a
```

```
mkdir -p ~/k8s-multi-tenant/tenant-b
```

```
cd ~/k8s-multi-tenant
```

Step 5: Create Deployment and Service for Tenant A

Create tenant-a-app.yaml in the tenant-a/ directory with the following contents:

```
apiVersion: apps/v1
```

```
kind: Deployment
metadata:
  name: tenant-a-app
  namespace: tenant-a
spec:
  replicas: 2
  selector:
    matchLabels:
      app: tenant-a-app
  template:
    metadata:
      labels:
        app: tenant-a-app
    spec:
      containers:
        - name: tenant-a-app
          image: nginx
```

```
apiVersion: v1
kind: Service
metadata:
  name: tenant-a-service
  namespace: tenant-a
spec:
  selector:
    app: tenant-a-app
  ports:
    - protocol: TCP
      port: 80
```

targetPort: 80

Apply the configuration:

```
kubectl apply -f tenant-a/tenant-a-app.yaml
```

Step 6: Restrict Network Access for Tenant A

Create tenant-a-restrict.yaml in the tenant-a/ directory with the following contents:

```
apiVersion: networking.k8s.io/v1
```

```
kind: NetworkPolicy
```

```
metadata:
```

```
  name: tenant-a-restrict
```

```
  namespace: tenant-a
```

```
spec:
```

```
  podSelector:
```

```
    matchLabels:
```

```
      app: tenant-a-app
```

```
  policyTypes:
```

```
  - Ingress
```

```
  ingress:
```

```
  - from:
```

```
    - podSelector:
```

```
      matchLabels:
```

```
        app: tenant-a-app
```

Apply the network policy:

```
kubectl apply -f tenant-a/tenant-a-restrict.yaml
```

Step 7: Create Deployment and Service for Tenant B

Create tenant-b-app.yaml in the tenant-b/ directory with the following contents:

```
apiVersion: apps/v1
```

```
kind: Deployment
metadata:
  name: tenant-b-app
  namespace: tenant-b
spec:
  replicas: 2
  selector:
    matchLabels:
      app: tenant-b-app
  template:
    metadata:
      labels:
        app: tenant-b-app
    spec:
      containers:
        - name: tenant-b-app
          image: nginx
```

```
apiVersion: v1
kind: Service
metadata:
  name: tenant-b-service
  namespace: tenant-b
spec:
  selector:
    app: tenant-b-app
  ports:
    - protocol: TCP
      port: 80
```

targetPort: 80

Apply the deployment:

```
kubectl apply -f tenant-b/tenant-b-app.yaml
```

Verify the deployment:

```
kubectl get pods -n tenant-b
```

```
kubectl get svc -n tenant-b
```

Step 8: Restrict Network Access for Tenant B

Create tenant-b-restrict.yaml in the tenant-b/ directory with the following contents:

```
apiVersion: networking.k8s.io/v1
```

```
kind: NetworkPolicy
```

```
metadata:
```

```
  name: tenant-b-restrict
```

```
  namespace: tenant-b
```

```
spec:
```

```
  podSelector:
```

```
    matchLabels:
```

```
      app: tenant-b-app
```

```
  policyTypes:
```

```
  - Ingress
```

```
  ingress:
```

```
  - from:
```

```
    - podSelector:
```

```
      matchLabels:
```

```
        app: tenant-b-app
```

Apply the network policy:

```
kubectl apply -f tenant-b/tenant-b-restrict.yaml
```

Step 9: Verify Network Policy

To verify the network policy for Tenant B, run the following commands:

```
kubectl get networkpolicy -n tenant-b
```

```
kubectl describe networkpolicy tenant-b-restrict -n tenant-b
```

Step 10: Final Folder Structure

The final folder structure should look like this:

```
k8s-multi-tenant/
```

```
|— tenant-a/
```

```
| |— tenant-a-app.yaml
```

```
| |— tenant-a-restrict.yaml
```

```
|— tenant-b/
```

```
| |— tenant-b-app.yaml
```

```
| |— tenant-b-restrict.yaml
```

Step 11: Test Tenant Isolation

Create a test pod in tenant-b and check access to tenant-a:

In worker docker run

```
docker pull alpine
```

```
kubectl run test-pod --image=alpine -n tenant-b --restart=Never -- sleep 3600
```

```
master@master-vn:~/k8s-multi-tenant$ kubectl run test-pod --image=alpine -n tenant-b --restart=Never -- sleep 3600
pod/test-pod created
```

```
kubectl exec -it test-pod -n tenant-b -- wget --spider tenant-a-service.tenant-a
```

```
master@master-vm:~$ kubectl run test-pod --image=alpine -n tenant-b --restart=Never -- sleep 3600
Error from server (AlreadyExists): pods "test-pod" already exists
master@master-vm:~$ kubectl exec -it test-pod -n tenant-b -- wget --spider tenant-a-service tenant-a
```

```
master@master-vm:~/k8s-multi-tenant$ kubectl get pod -n tenant-a -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GAT
tenant-a-app-57856ccbdcc-44552	1/1	Running	2 (3h56m ago)	2d	192.168.94.199	worker1-vm	<none>		<none>	
tenant-a-app-57856ccbdcc-cwpnp	1/1	Running	2 (3h55m ago)	2d	192.168.114.72	worker2-vm	<none>		<none>	

```
master@master-vm:~/k8s-multi-tenant$ kubectl get pod -n tenant-b -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GA
tenant-b-app-bbb987489-8bj55	1/1	Running	2 (3h55m ago)	2d	192.168.114.71	worker2-vm	<none>		<none>	
tenant-b-app-bbb987489-hzcnk	1/1	Running	2 (3h56m ago)	2d	192.168.94.201	worker1-vm	<none>		<none>	
test-pod	1/1	Running	0	7m26s	192.168.114.74	worker2-vm	<none>		<none>	

```
master@master-vm:~/k8s-multi-tenant$
```